
MODULE *ZeusOwnershipMeta*

EXTENDS *Integers, FiniteSets*

CONSTANTS *LB_NODES* and *APP_NODES* must not intersect and neither should contain 0.

LB_NODES,
APP_NODES,
O_MAX_VERSION,
O_MAX_FAILURES,
O_MAX_DATA_VERSION

VARIABLES variable prefixes – \rightarrow *o*: ownership, *r*: request, *t*: transactional, *m*: membership

VECTORS indexed by *node_id*

oTS,
oState,
oDriver,
oVector,
oRcvACKs,

No readers/owner: *.readers* = {} / *.owner* = 0

rTS,
rID,
rType,
rEID,

since we do not have message loss timeouts we use this to
track epoch of last issued *INV*s for replays

tState,
tVersion,
tRcvACKs,

tVesion suffice to represent *tData* | = 0 – \rightarrow no data | > 0 data (reader / owner)

GLOBAL variables

oMsgs,

mAliveNodes, membership

mEID, membership epoch *id*

committedREQs, only to check invariant that exactly one of concurrent *REQs* is committed

committedRTS only to emulate *FIFO REQ* channels (*i.e.*, do not re execute same client requests)

$vars \triangleq \langle oTS, oState, oDriver, oVector, oRcvACKs, rTS, rID, rType, rEID,$
 $tState, tVersion, tRcvACKs, oMsgs, mAliveNodes, mEID, committedREQs, committedRTS \rangle$

Helper operators

$O_NODES \triangleq LB_NODES \cup APP_NODES$

$O_NODES_0 \triangleq O_NODES \cup \{0\}$

$LB_NODES_0 \triangleq LB_NODES \cup \{0\}$

$APP_NODES_0 \triangleq APP_NODES \cup \{0\}$

$LB_LIVE_NODES \triangleq LB_NODES \cap mAliveNodes$

$APP_LIVE_NODES \triangleq APP_NODES \cap mAliveNodes$

$LB_LIVE_ARBITERS(driver) \triangleq LB_LIVE_NODES \setminus \{driver\}$ all arbiters except driver and owner

ASSUME $LB_NODES \cap APP_NODES = \{\}$
 ASSUME $\forall k \in O_NODES : k \neq 0$ we use 0 as the default noop

Useful Unchanged shortcuts

$unchanged_M \triangleq \text{UNCHANGED } \langle oMsgs \rangle$
 $unchanged_m \triangleq \text{UNCHANGED } \langle mEID, mAliveNodes \rangle$
 $unchanged_t \triangleq \text{UNCHANGED } \langle tState, tVersion, tRcvACKs \rangle$
 $unchanged_r \triangleq \text{UNCHANGED } \langle rID, rTS, rEID, rType \rangle$
 $unchanged_c \triangleq \text{UNCHANGED } \langle committedREQs, committedRTS \rangle$
 $unchanged_o \triangleq \text{UNCHANGED } \langle oState, oDriver, oVector, oRcvACKs, oTS \rangle$
 $unchanged_mc \triangleq unchanged_m \wedge unchanged_c$
 $unchanged_mtc \triangleq unchanged_mc \wedge unchanged_t$
 $unchanged_mtr \triangleq unchanged_m \wedge unchanged_t \wedge unchanged_r$
 $unchanged_Mrc \triangleq unchanged_r \wedge unchanged_c \wedge unchanged_M$
 $unchanged_mrco \triangleq unchanged_mc \wedge unchanged_r \wedge unchanged_o$
 $unchanged_mtco \triangleq unchanged_mtc \wedge unchanged_o$
 $unchanged_mtrc \triangleq unchanged_mtc \wedge unchanged_r$
 $unchanged_Mtrc \triangleq unchanged_Mrc \wedge unchanged_t$
 $unchanged_Mmrc \triangleq unchanged_Mrc \wedge unchanged_m$
 $unchanged_mtrco \triangleq unchanged_mtrc \wedge unchanged_o$
 $unchanged_Mmrco \triangleq unchanged_mrco \wedge unchanged_M$
 $unchanged_Mmtrc \triangleq unchanged_mtrc \wedge unchanged_M$

Type definitions

$Type_oTS \triangleq [ver : 0 \dots O_MAX_VERSION, tb : LB_NODES_0]$
 $Type_rTS \triangleq [ver : 0 \dots O_MAX_VERSION, tb : APP_NODES_0]$
 $Type_tState \triangleq \{ \text{"valid"}, \text{"invalid"}, \text{"write"} \}$ readers can be in valid and invalid and owner in valid and write
 $Type_oState \triangleq \{ \text{"valid"}, \text{"invalid"}, \text{"drive"}, \text{"request"} \}$ all nodes start from valid
 $Type_rType \triangleq \{ \text{"add-owner"}, \text{"change-owner"}, \text{"add-reader"}, \text{"rm-reader"}, \text{"NOOP"} \}$
 $Type_oVector \triangleq [readers : \text{SUBSET } APP_NODES, owner : APP_NODES_0]$

$Type_oMessage \triangleq$ Msgs exchanged by the sharding protocol
 $[type : \{ \text{"REQ"} \},$
 $\begin{array}{ll} rTS & : Type_rTS, \\ rID & : Nat, \\ rType & : Type_rType, \\ epochID & : 0 \dots O_MAX_FAILURES \end{array}$
 \cup
 $[type : \{ \text{"NACK"} \},$
 $\begin{array}{ll} rTS & : Type_rTS, \\ rID & : Nat \end{array}$
 \cup
 $[type : \{ \text{"S_INV"} \},$
 $\begin{array}{ll} sender & : O_NODES, \\ driver & : O_NODES, \\ rTS & : Type_rTS, \end{array}$

$$\begin{aligned}
& rID : Nat, \\
& oTS : Type_oTS, \\
& oVector : Type_oVector, \\
& rType : Type_rType, \\
& epochID : 0 \dots O_MAX_FAILURES] \\
\cup \\
& [type : \{ "S_ACK" \}, \quad sender : O_NODES, \\
& \quad oTS : Type_oTS, \\
& \quad tVersion : 0 \dots O_MAX_DATA_VERSION, \quad \text{emulates data send as well} \\
& \quad epochID : 0 \dots O_MAX_FAILURES] \\
\cup \\
& [type : \{ "RESP" \}, \quad oVector : Type_oVector, \\
& \quad oTS : Type_oTS, \\
& \quad rTS : Type_rTS, \\
& \quad preOwner, \setminus * \text{ pre-request owner is not needed for model check (since we model bcast messages)} \\
& \quad tVersion : 0 \dots O_MAX_DATA_VERSION, \\
& \quad epochID : 0 \dots O_MAX_FAILURES] \\
\cup \\
& [type : \{ "S_VAL" \}, \quad oTS : Type_oTS, \\
& \quad epochID : 0 \dots O_MAX_FAILURES] \\
\\
Type_tMessage \triangleq \quad \text{msgs exchanged by the transactional reliable commit protocol} \\
& [type : \{ "T_INV", "T_ACK", "T_VAL" \}, \quad tVersion : Nat, \\
& \quad sender : O_NODES, \\
& \quad epochID : 0 \dots O_MAX_FAILURES]
\end{aligned}$$

Type check and initialization

$$\begin{aligned}
OTypeOK \triangleq \quad & \text{The type correctness invariant} \\
& \wedge oTS \in [O_NODES \rightarrow Type_oTS] \\
& \wedge oState \in [O_NODES \rightarrow Type_oState] \\
& \wedge oDriver \in [O_NODES \rightarrow O_NODES_0] \\
& \wedge oVector \in [O_NODES \rightarrow Type_oVector] \\
& \wedge \forall n \in O_NODES : oRcvACKs[n] \subseteq (O_NODES \setminus \{n\}) \\
& \wedge rTS \in [O_NODES \rightarrow Type_rTS] \\
& \wedge rID \in [O_NODES \rightarrow 0 \dots O_MAX_VERSION] \\
& \wedge rType \in [O_NODES \rightarrow Type_rType] \\
& \wedge rEID \in [O_NODES \rightarrow 0 \dots (Cardinality(O_NODES) - 1)] \\
& \wedge tVersion \in [O_NODES \rightarrow 0 \dots O_MAX_DATA_VERSION] \\
& \wedge tState \in [O_NODES \rightarrow Type_tState] \\
& \wedge \forall n \in O_NODES : tRcvACKs[n] \subseteq (O_NODES \setminus \{n\}) \\
& \wedge committedREQs \subseteq Type_oTS \\
& \wedge committedRTS \subseteq Type_rTS
\end{aligned}$$

$$\begin{aligned}
\wedge oMsgs &\subseteq (Type_oMessage \cup Type_tMessage) \\
\wedge mEID &\in 0 \dots (Cardinality(O_NODES) - 1) \\
\wedge mAliveNodes &\subseteq O_NODES
\end{aligned}$$

$OInit \triangleq$ The initial predicate

$$\begin{aligned}
\wedge oTS &= [n \in O_NODES \mapsto [ver \mapsto 0, tb \mapsto 0]] \\
\wedge oState &= [n \in O_NODES \mapsto \text{"valid"}] \\
\wedge oDriver &= [n \in O_NODES \mapsto 0] \\
\wedge oVector &= [n \in O_NODES \mapsto [readers \mapsto \{\}, owner \mapsto 0]] \\
\wedge oRcvACKs &= [n \in O_NODES \mapsto \{\}] \\
\wedge rTS &= [n \in O_NODES \mapsto [ver \mapsto 0, tb \mapsto 0]] \\
\wedge rID &= [n \in O_NODES \mapsto 0] \\
\wedge rEID &= [n \in O_NODES \mapsto 0] \\
\wedge rType &= [n \in O_NODES \mapsto \text{"NOOP"}] \\
\wedge tVersion &= [n \in O_NODES \mapsto 0] \\
\wedge tState &= [n \in O_NODES \mapsto \text{"valid"}] \\
\wedge tRcvACKs &= [n \in O_NODES \mapsto \{\}] \\
\wedge committedRTS &= \{\} \\
\wedge committedREQs &= \{\} \\
\wedge oMsgs &= \{\} \\
\wedge mEID &= 0 \\
\wedge mAliveNodes &= O_NODES
\end{aligned}$$

$$\begin{aligned}
Min(S) &\triangleq \text{CHOOSE } x \in S : \forall y \in S \setminus \{x\} : y > x \\
set_wo_min(S) &\triangleq S \setminus \{Min(S)\}
\end{aligned}$$

First Command executed once after $OInit$ to initialize owner/readers and $oVector$ state

$$\begin{aligned}
OInit_min_owner_rest_readers &\triangleq \\
&\wedge \forall x \in O_NODES : tVersion[x] = 0 \\
&\wedge tVersion' = [n \in O_NODES \mapsto \text{IF } n \in LB_NODES \text{ THEN } 0 \text{ ELSE } 1] \\
&\wedge oVector' = [n \in O_NODES \mapsto \text{IF } n \in set_wo_min(APP_NODES) \\
&\quad \text{THEN } oVector[n] \\
&\quad \text{ELSE } [readers \mapsto set_wo_min(APP_NODES), \\
&\quad \quad owner \mapsto Min(APP_NODES)]] \\
&\wedge unchanged_Mmrc \\
&\wedge UNCHANGED \langle tState, oState, oDriver, oRcvACKs, oTS, tRcvACKs \rangle
\end{aligned}$$

Helper functions

$$\begin{aligned}
has_data(n) &\triangleq tVersion[n] > 0 \\
has_valid_data(n) &\triangleq \wedge has_data(n) \\
&\quad \wedge tState[n] = \text{"valid"} \\
is_owner(n) &\triangleq \wedge has_data(n) \\
&\quad \wedge oVector[n].owner = n \\
is_valid_owner(n) &\triangleq \wedge is_owner(n)
\end{aligned}$$

$$\begin{aligned}
& \wedge oState[n] = \text{"valid"} \\
is_reader(n) & \triangleq \begin{aligned} & \wedge has_data(n) \\ & \wedge \neg is_owner(n) \\ & \wedge n \notin LB_NODES \end{aligned} \\
is_live_arbiter(n) & \triangleq \begin{aligned} & \vee n \in LB_LIVE_NODES \\ & \vee is_owner(n) \end{aligned} \\
is_valid_live_arbiter(n) & \triangleq \begin{aligned} & \wedge is_live_arbiter(n) \\ & \wedge oState[n] = \text{"valid"} \end{aligned} \\
is_requester(n) & \triangleq \begin{aligned} & \wedge n \in APP_LIVE_NODES \\ & \wedge \neg is_owner(n) \end{aligned} \\
is_valid_requester(n) & \triangleq \begin{aligned} & \wedge is_requester(n) \\ & \wedge oState[n] = \text{"valid"} \end{aligned} \\
is_in_progress_requester(n) & \triangleq \begin{aligned} & \wedge is_requester(n) \\ & \wedge oState[n] = \text{"request"} \end{aligned} \\
requester_is_alive(n) & \triangleq rTS[n].tb \in mAliveNodes
\end{aligned}$$

Timestamp Comparison Helper functions

$$\begin{aligned}
is_equalTS(ts1, ts2) & \triangleq \begin{aligned} & \wedge ts1.ver = ts2.ver \\ & \wedge ts1.tb = ts2.tb \end{aligned} \\
is_greaterTS(ts1, ts2) & \triangleq \begin{aligned} & \vee ts1.ver > ts2.ver \\ & \vee \wedge ts1.ver = ts2.ver \\ & \quad \wedge ts1.tb > ts2.tb \end{aligned} \\
is_greaterreqTS(ts1, ts2) & \triangleq \begin{aligned} & \vee is_equalTS(ts1, ts2) \\ & \vee is_greaterTS(ts1, ts2) \end{aligned} \\
is_smallerTS(ts1, ts2) & \triangleq \neg is_greaterreqTS(ts1, ts2)
\end{aligned}$$

Request type Helper functions

$$is_non_sharing_req(n) \triangleq (rType[n] = \text{"add-owner"} \vee rType[n] = \text{"add-reader"})$$

Post o_vector based on request type and r (requester or 0 if requester is not alive)

$$post_oVec(n, r, pre_oVec) \triangleq$$

```

IF ( $rType[n] = \text{"add-owner"} \vee rType[n] = \text{"change-owner"}$ )
  THEN [ $owner \mapsto r$ ,
         $readers \mapsto (pre\_oVec.readers \cup \{pre\_oVec.owner\}) \setminus \{r, 0\}$ ]
ELSE [ $owner \mapsto pre\_oVec.owner$ ,
       $readers \mapsto$  IF  $rType[n] = \text{"remove-reader"}$ 
        THEN  $pre\_oVec.readers \setminus \{r, 0\}$ 
        ELSE  $rType[n] = \text{"add-reader"}$ 
          ( $pre\_oVec.readers \cup \{r\} \setminus \{0\}$ )

```

Message Helper functions

Used only to emulate *FIFO REQ* channels (and not re-execute already completed *REQs*)
 $not_completed_rTS(r_ts) \triangleq \forall c_rTS \in committedRTS : c_rTS \neq r_ts$

Messages in *oMsgs* are only appended to this variable (not removed once delivered)
 intentionally to check protocols tolerance in duplicates and reorderings
 $send_omsg(m) \triangleq oMsgs' = oMsgs \cup \{m\}$

$o_send_req(r_ts, r_id, r_type) \triangleq$
 $send_omsg([type \mapsto \text{"REQ"},$
 $rTS \mapsto r_ts,$
 $rID \mapsto r_id,$
 $rType \mapsto r_type,$
 $epochID \mapsto mEID])$

$o_send_nack(r_ts, r_id) \triangleq$
 $send_omsg([type \mapsto \text{"NACK"},$
 $rTS \mapsto r_ts,$
 $rID \mapsto r_id])$

$o_send_inv(sender, driver, o_ts, o_vec, r_ts, r_id, r_type) \triangleq$
 $send_omsg([type \mapsto \text{"S_INV"},$
 $sender \mapsto sender,$
 $driver \mapsto driver,$
 $oTS \mapsto o_ts,$
 $oVector \mapsto o_vec,$
 $rTS \mapsto r_ts,$
 $rID \mapsto r_id,$
 $rType \mapsto r_type,$
 $epochID \mapsto mEID])$

$o_send_ack(sender, o_ts, t_version) \triangleq$
 $send_omsg([type \mapsto \text{"S_ACK"},$
 $sender \mapsto sender,$
 $oTS \mapsto o_ts,$

$$\begin{aligned}
& tVersion \mapsto t_version, \\
& epochID \mapsto mEID]) \\
o_send_resp(r_ts, o_ts, o_vec, t_version) & \triangleq \\
& send_omsg([type \mapsto \text{"RESP"}, \\
& \quad oVector \mapsto o_vec, \\
& \quad oTS \mapsto o_ts, \\
& \quad rTS \mapsto r_ts, \\
& \quad tVersion \mapsto t_version, \\
& \quad epochID \mapsto mEID]) \\
o_send_val(o_ts) & \triangleq \\
& send_omsg([type \mapsto \text{"S_VAL"}, \\
& \quad oTS \mapsto o_ts, \\
& \quad epochID \mapsto mEID \quad]])
\end{aligned}$$

Operators to check received messages (m stands for message)

$$\begin{aligned}
o_rcv_req(m) & \triangleq \\
& \wedge m.type = \text{"REQ"} \\
& \wedge m.epochID = mEID \\
& \wedge not_completed_rTS(m.rTS) \\
o_rcv_nack(m, receiver) & \triangleq \\
& \wedge m.type = \text{"NACK"} \\
& \wedge m.rTS = rTS[receiver] \\
& \wedge m.rID = rID[receiver] \\
o_rcv_resp(m, receiver) & \triangleq \\
& \wedge m.type = \text{"RESP"} \\
& \wedge m.epochID = mEID \\
& \wedge m.rTS = rTS[receiver] \\
o_rcv_inv(m, receiver) & \triangleq \\
& \wedge m.type = \text{"S_INV"} \\
& \wedge m.epochID = mEID \\
& \wedge m.sender \neq receiver \\
o_rcv_inv_equal_ts(m, receiver) & \triangleq \\
& \wedge o_rcv_inv(m, receiver) \\
& \wedge is_equalTS(m.oTS, oTS[receiver]) \\
o_rcv_inv_smaller_ts(m, receiver) & \triangleq \\
& \wedge o_rcv_inv(m, receiver) \\
& \wedge is_smallerTS(m.oTS, oTS[receiver]) \\
o_rcv_inv_greater_ts(m, receiver) & \triangleq \\
& \wedge o_rcv_inv(m, receiver) \\
& \wedge is_greaterTS(m.oTS, oTS[receiver])
\end{aligned}$$

$$\begin{aligned}
o_rcv_inv_greaterreq_ts(m, receiver) &\triangleq \\
&\wedge o_rcv_inv(m, receiver) \\
&\wedge \neg is_smallerTS(m.oTS, oTS[receiver])
\end{aligned}$$

$$\begin{aligned}
o_rcv_ack(m, receiver) &\triangleq \\
&\wedge m.type = \text{"S_ACK"} \\
&\wedge m.epochID = mEID \\
&\wedge m.sender \neq receiver \\
&\wedge oState[receiver] = \text{"drive"} \\
&\wedge m.sender \notin oRcvACKs[receiver] \\
&\wedge is_equalTS(m.oTS, oTS[receiver])
\end{aligned}$$

$$\begin{aligned}
o_rcv_val(m, receiver) &\triangleq \\
&\wedge m.type = \text{"S_VAL"} \\
&\wedge m.epochID = mEID \\
&\wedge oState[receiver] \neq \text{"valid"} \\
&\wedge is_equalTS(m.oTS, oTS[receiver])
\end{aligned}$$

Used to not re-issue messages that already exists (and bound the state space)

$$\begin{aligned}
msg_not_exists(o_rcv_msg(-, -), receiver) &\triangleq \\
&\neg \exists mm \in oMsgs : o_rcv_msg(mm, receiver)
\end{aligned}$$

$$rcved_acks_from_set(n, set) \triangleq set \subseteq oRcvACKs[n]$$

Check if all acknowledgments from arbiters have been received

$$\begin{aligned}
has_rcved_all_ACKs(n) &\triangleq \\
&\wedge rEID[n] = mEID \\
&\wedge \text{IF } oVector[n].owner \neq 0 \\
&\quad \text{THEN } rcved_acks_from_set(n, \{oVector[n].owner\} \cup LB_LIVE_ARBITERS(n)) \\
&\quad \text{ELSE } \vee \wedge \neg requester_is_alive(n) \\
&\quad \quad \wedge rcved_acks_from_set(n, LB_LIVE_ARBITERS(n)) \\
&\quad \vee \wedge oVector[n].readers \neq \{\} \\
&\quad \quad \wedge \exists x \in oVector[n].readers : rcved_acks_from_set(n, \{x\} \cup LB_LIVE_ARBITERS(n))
\end{aligned}$$

message helper functions related to transactions

$$\begin{aligned}
t_send(n, msg_type, t_ver) &\triangleq \\
&\quad send_omsg([type \mapsto msg_type, \\
&\quad \quad tVersion \mapsto t_ver, \\
&\quad \quad sender \mapsto n, \\
&\quad \quad epochID \mapsto mEID \quad])
\end{aligned}$$

$$\begin{aligned}
t_rcv_inv(m, receiver) &\triangleq \\
&\wedge m.type = \text{"T_INV"} \\
&\wedge m.epochID = mEID \\
&\wedge m.sender \neq receiver
\end{aligned}$$

$$\begin{aligned}
t_rcv_ack(m, receiver) &\triangleq \\
&\wedge m.type = \text{"T_ACK"} \\
&\wedge m.epochID = mEID \\
&\wedge m.sender \neq receiver \\
&\wedge tState[receiver] = \text{"write"} \\
&\wedge m.sender \notin tRcvACKs[receiver] \\
&\wedge m.tVersion = tVersion[receiver]
\end{aligned}$$

$$\begin{aligned}
t_rcv_val(m, receiver) &\triangleq \\
&\wedge m.type = \text{"T_VAL"} \\
&\wedge m.epochID = mEID \\
&\wedge tState[receiver] \neq \text{"valid"} \\
&\wedge m.tVersion = tVersion[receiver]
\end{aligned}$$

Protocol Invariants:

Valid data are consistent

$$\begin{aligned}
CONSISTENT_DATA &\triangleq \\
&\forall k, n \in APP_LIVE_NODES : \vee \neg has_valid_data(k) \\
&\quad \vee \neg has_valid_data(n) \\
&\quad \vee tVersion[n] = tVersion[k]
\end{aligned}$$

Amongst concurrent sharing requests only one succeeds

The invariant that we cannot have two *REQs* committed with same versions
(i.e., that read and modified the same sharing vector)

$$\begin{aligned}
ONLY_ONE_CONC_REQ_COMMITTS &\triangleq \\
&\forall x, y \in committedREQs : \vee x.ver \neq y.ver \\
&\quad \vee x.tb = y.tb
\end{aligned}$$

There is always at most one valid owner

$$\begin{aligned}
AT_MOST_ONE_OWNER &\triangleq \\
&\forall n, m \in mAliveNodes : \vee \neg is_valid_owner(n) \\
&\quad \vee \neg is_valid_owner(m) \\
&\quad \vee m = n
\end{aligned}$$

Valid owner has the most up-to-date data and version among live replicas

$$\begin{aligned}
OWNER_LATEST_DATA &\triangleq \\
&\forall o, k \in mAliveNodes : \vee \neg is_valid_owner(o) \\
&\quad \vee \neg has_data(o) \\
&\quad \vee tVersion[o] \geq tVersion[k]
\end{aligned}$$

All valid sharers (*LB* + *owner*) agree on their sharing vectors (and *TS*)

$$\begin{aligned}
CONSISTENT_SHARERS &\triangleq \\
&\forall k, n \in mAliveNodes : \vee \neg is_valid_live_arbiter(n) \\
&\quad \vee \neg is_valid_live_arbiter(k) \\
&\quad \vee \wedge oTS[n] = oTS[k]
\end{aligned}$$

$$\wedge oVector[n] = oVector[k]$$

$$\begin{aligned} CONSISTENT_OVECTORS_Fwd &\triangleq \\ \forall n \in mAliveNodes : &\vee \neg is_valid_live_arbiter(n) \\ &\vee \wedge \forall r \in oVector[n].readers : \\ &\quad \wedge has_data(r) \\ &\quad \wedge \neg is_valid_owner(r) \\ &\wedge \vee oVector[n].owner = 0 \\ &\quad \vee is_owner(oVector[n].owner) \end{aligned}$$

$$\begin{aligned} CONSISTENT_OVECTORS_Reverse_owner &\triangleq \\ \forall o, n \in mAliveNodes : &\vee \neg is_valid_owner(o) \\ &\vee \neg is_valid_live_arbiter(n) \\ &\vee oVector[n].owner = o \end{aligned}$$

$$\begin{aligned} CONSISTENT_OVECTORS_Reverse_readers &\triangleq \\ \forall r, n \in mAliveNodes : &\vee \neg is_reader(r) \\ &\vee \neg is_valid_live_arbiter(n) \\ &\vee r \in oVector[n].readers \end{aligned}$$

The owner and readers are always correctly reflected by any valid sharing vectors

$$\begin{aligned} CONSISTENT_OVECTORS &\triangleq \\ &\wedge CONSISTENT_OVECTORS_Fwd \\ &\wedge CONSISTENT_OVECTORS_Reverse_owner \\ &\wedge CONSISTENT_OVECTORS_Reverse_readers \end{aligned}$$